



# **CISC 867: PROJECT 1**

**Leaf Classification**

---

**Lomai Mohamed Saadeldin**

**ID: 20398049**

---

**Supervised by: Prof. Hazem Abbas**



## ✓ Problem Formulation:

### ▪ Data Description:

- The dataset consists approximately 1,584 images of leaf specimens (16 samples each of 99 species) which have been converted to binary black leaves against white backgrounds.
- Three sets of features are also provided per image: a shape contiguous descriptor, an interior texture histogram, and a fine-scale margin histogram. For each feature, a 64-attribute vector is given per leaf sample.
- Note that of the original 100 species, we have eliminated one on account of incomplete associated data in the original dataset.

### ▪ The Problems:

- We will use this data using a neural network architecture.

### ▪ Input:

- Features collected from numbers of species of plant.

### ▪ Output:

- Predicted species for leaves.

### ▪ Challenges :

- Clean the data set.
- Try different hyperparameters for the network.

### ▪ Impact :

- Predicting the species of the leaf that will lead to a successful match.

## Part I: Data Preparation

1) import modules to dealing with data set

```
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import numpy as np # linear algebra
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn import preprocessing
from sklearn.preprocessing import StandardScaler, OneHotEncoder

from matplotlib.pyplot import figure, imshow, axis

#DEEP LEARNING
from keras.layers.core import Dense
from keras.layers import Dropout
from tensorflow import keras
from keras.models import Sequential
from keras import regularizers
import tensorflow as tf
from tensorflow import keras
from keras import layers
from keras.optimizers import Optimizer
from keras.optimizers import SGD
from keras.optimizers import Adam
from keras.optimizers import RMSprop
```

2) Download the data file, load it:


```
[4] #read dataset into files
train=pd.read_csv('/content/train.csv') #train dataset
test=pd.read_csv('/content/test.csv') #test dataset
```

### 3) Describe the data

```
[9] train.describe()
```

	id	margin1	margin2	margin3	margin4	margin5	margin6	margin7	margin
<b>count</b>	990.000000	990.000000	990.000000	990.000000	990.000000	990.000000	990.000000	990.000000	990.000000
<b>mean</b>	799.595960	0.017412	0.028539	0.031988	0.023280	0.014264	0.038579	0.019202	0.00108
<b>std</b>	452.477568	0.019739	0.038855	0.025847	0.028411	0.018390	0.052030	0.017511	0.00274
<b>min</b>	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
<b>25%</b>	415.250000	0.001953	0.001953	0.013672	0.005859	0.001953	0.000000	0.005859	0.000000
<b>50%</b>	802.500000	0.009766	0.011719	0.025391	0.013672	0.007812	0.015625	0.015625	0.000000
<b>75%</b>	1195.500000	0.025391	0.041016	0.044922	0.029297	0.017578	0.056153	0.029297	0.000000
<b>max</b>	1584.000000	0.087891	0.205080	0.156250	0.169920	0.111330	0.310550	0.091797	0.03125

8 rows × 10 columns



#### 4) Clean the data set

check the null value

```
[13] #check the null values (train set)
train.isnull().sum()
```

```
id          0
species     0
margin1     0
margin2     0
margin3     0
..
texture60   0
texture61   0
texture62   0
texture63   0
texture64   0
Length: 194, dtype: int64
```

```
[14] #check the null values (test set)
test.isnull().sum()
```

```
id          0
margin1     0
margin2     0
margin3     0
margin4     0
..
texture60   0
texture61   0
texture62   0
texture63   0
texture64   0
Length: 193, dtype: int64
```

Check for duplicates

```
[15]
#check for duplicates (train set)
train.duplicated().sum()
```

```
0
```

```
[16] #check for duplicates (test set)
test.duplicated().sum()
```

```
0
```


## 5) Draw some of the images

```
[17] import cv2
import glob


images = []
files = glob.glob ("/content/drive/MyDrive/images/*.jpg")
for myFile in files:
    image = cv2.imread(myFile, cv2.IMREAD_GRAYSCALE)
    images.append(image)


[31] plt.figure(figsize=(15,10))
plt.imshow(images[0],cmap='gray', interpolation='nearest')
```

<matplotlib.image.AxesImage at 0x7f0038cb4b20>




```
plt.figure(figsize=(15,10))  
plt.imshow(images[2],cmap='gray', interpolation='nearest')
```

 <matplotlib.image.AxesImage at 0x7f0034a45070>



```
plt.figure(figsize=(15,10))  
plt.imshow(images[5],cmap='gray', interpolation='nearest')
```

<matplotlib.image.AxesImage at 0x7f00341bdfd0>



↑ ↓ 🔗 💬 ✎ 📄 🗑️ ⋮



## 6) correlation analysis

```
correlation_matrix = train.corr().abs()  
print(correlation_matrix)
```

```
id      id      margin1  margin2  margin3  margin4  margin5  \  
id      1.000000  0.011673  0.027565  0.059533  0.001639  0.002419  \  
margin1 0.011673  1.000000  0.806390  0.182829  0.297807  0.475874  \  
margin2 0.027565  0.806390  1.000000  0.204640  0.315953  0.444312  \  
margin3 0.059533  0.182829  0.204640  1.000000  0.120042  0.185007  \  
margin4 0.001639  0.297807  0.315953  0.120042  1.000000  0.029480  \  
...      ...      ...      ...      ...      ...      ...      \  
texture60 0.000823  0.035072  0.081069  0.019850  0.052317  0.006542  \  
texture61 0.026319  0.007581  0.007057  0.084957  0.320644  0.109229  \  
texture62 0.032873  0.033159  0.037405  0.081999  0.073886  0.151675  \  
texture63 0.024299  0.075171  0.098957  0.148193  0.050970  0.022299  \  
texture64 0.035396  0.030414  0.029532  0.061780  0.014343  0.148834  \  
  
id      margin6  margin7  margin8  margin9  ...  texture55  texture56  \  
id      0.051818  0.061214  0.039509  0.070954  ...  0.040292  0.005132  \  
margin1 0.767718  0.066273  0.094137  0.181496  ...  0.137158  0.047771  \  
margin2 0.825762  0.083273  0.086428  0.120276  ...  0.154407  0.021096  \  
margin3 0.163976  0.095449  0.024350  0.000042  ...  0.047347  0.027618  \  
margin4 0.261437  0.268271  0.047693  0.227543  ...  0.071974  0.009537  \  
...      ...      ...      ...      ...  ...  ...      ...      \  
texture60 0.066262  0.034094  0.048647  0.028292  ...  0.129365  0.004412  \  
texture61 0.050498  0.163375  0.079283  0.088517  ...  0.002235  0.053707  \  
texture62 0.031555  0.015391  0.048843  0.031954  ...  0.217239  0.171577  \  
texture63 0.132087  0.001364  0.027758  0.119494  ...  0.207887  0.002057  \  
texture64 0.003164  0.068512  0.003191  0.097760  ...  0.095205  0.095913  \  
  
id      texture57  texture58  texture59  texture60  texture61  texture62  \  
id      0.043101  0.063337  0.007915  0.000823  0.026319  0.032873  \  
margin1 0.126227  0.024139  0.168201  0.035072  0.007581  0.033159  \  
margin2 0.123834  0.063654  0.157842  0.081069  0.007057  0.037405  \  
margin3 0.007261  0.021390  0.033505  0.019850  0.084957  0.081999  \  
margin4 0.050529  0.044318  0.088857  0.052317  0.320644  0.073886  \  
...      ...      ...      ...      ...      ...      ...      \  
texture60 0.155187  0.240704  0.183369  1.000000  0.051838  0.265879  \  
texture61 0.072814  0.084638  0.023539  0.051838  1.000000  0.063582  \  
texture62 0.283316  0.563088  0.128010  0.265879  0.063582  1.000000  \  
texture63 0.064724  0.059866  0.156568  0.089679  0.068065  0.058189  \  
texture64 0.224686  0.269157  0.015374  0.190194  0.036374  0.245527  \
```



[ ]



## 7) split data into train set and test set

```
[ ] x = train.drop('species',axis = 1)
    Y = train.species.values
    from sklearn.preprocessing import LabelEncoder
    label_encoder = LabelEncoder()
    y = label_encoder.fit_transform(Y)
```



```
x_train, x_test, y_train, y_test = train_test_split(x,y, random_state = 5, train_size=0.75,shuffle = True)
```

## 8) Normalize data



```
scaler = StandardScaler()  
scaler.fit(x_train)  
x_train=scaler.transform(x_train)  
x_test=scaler.transform(x_test)
```

```
[ ] print('mean',scaler.mean_)  
    print('scale',scaler.scale_)
```


```
mean [7.83587601e+02 1.70990108e-02 2.80518248e-02 3.21133841e-02  
2.28478140e-02 1.46404784e-02 3.82439313e-02 1.92074434e-02  
1.06863208e-03 7.43858221e-03 1.85072790e-02 2.38112318e-02  
1.20872049e-02 4.08340647e-02 8.00983962e-03 1.59171348e-02  
1.18447439e-04 1.49379515e-02 1.98286550e-02 1.22582911e-02  
1.30348140e-02 1.90626523e-02 6.16465903e-03 1.14501078e-03  
7.84665499e-03 8.85481806e-03 1.87363059e-02 5.69612129e-03  
1.57223464e-02 2.80860162e-02 1.63698666e-02 1.10843154e-02  
1.00314367e-02 1.94311792e-02 1.03967790e-03 1.30716779e-02  
1.79650202e-02 1.62725121e-02 3.12131307e-02 1.51880040e-02  
8.27043396e-03 1.14159771e-02 1.70016132e-02 1.87968167e-02  
1.28952965e-02 2.42060701e-02 1.00577183e-02 2.50510040e-02  
2.74858666e-02 8.82587197e-03 1.39060970e-02 2.55643208e-02  
2.94542318e-03 2.43956199e-02 1.02472790e-02 1.86888935e-02  
5.89615364e-03 1.26136429e-02 1.93548464e-02 3.06893005e-02  
1.20240323e-02 1.39500809e-03 5.15124394e-03 2.54853598e-02  
4.59581671e-03 7.28471941e-04 7.07037129e-04 6.82184960e-04  
6.59185903e-04 6.38773410e-04 6.21596954e-04 6.07364737e-04  
5.93877190e-04 5.80945175e-04 5.68772197e-04 5.59865303e-04  
5.51597615e-04 5.44224272e-04 5.39178863e-04 5.35114623e-04  
5.33399751e-04 5.32427930e-04 5.31567396e-04 5.31953889e-04  
5.36114877e-04 5.43444691e-04 5.52944600e-04 5.64752032e-04  
5.75690770e-04 5.88560433e-04 6.02043435e-04 6.15739018e-04  
6.31904035e-04 6.51519131e-04 6.75313616e-04 6.99679654e-04  
7.21895206e-04 7.26824232e-04 7.06407965e-04 6.79579040e-04  
6.54815082e-04 6.32504625e-04 6.11320988e-04 5.92783865e-04  
5.75329730e-04 5.61806011e-04 5.51197875e-04 5.44890580e-04  
5.38795937e-04 5.35181608e-04 5.33334003e-04 5.34240786e-04  
5.35375089e-04 5.37140028e-04 5.38120687e-04 5.37302439e-04  
5.38117682e-04 5.41949146e-04 5.46436186e-04 5.52663935e-04  
5.59650171e-04 5.68595879e-04 5.81757454e-04 5.97590139e-04  
6.16994602e-04 6.38255210e-04 6.64383302e-04 6.93275229e-04  
7.22135782e-04 2.16515323e-02 1.18214259e-02 1.01104582e-02  
1.54618491e-02 2.65199286e-02 1.01499272e-02 1.66345121e-02  
1.94891361e-02 1.49972237e-02 1.96405512e-02 1.88758544e-02  
0.71011010e-02 0.01160700e-02 1.07110010e-02 0.50101070e-02
```

## 9) Data Visualizing

Visualize the data using proper visualization methods.

### Sample of the species

```
[ ] Figure = plt.figure(figsize =(15, 10))  
plt.pie(y[:10],labels=train['species'][:10]);
```





## Part II: Training a neural network

### 1)Build ANN Model

```
[50]
def ANN(batch_size, hidden_dropout, Regularization, learning_rate, hidden_Size, n_epochs):
    #Build ANN MODEL
    Model = Sequential()
    Model.add(keras.layers.Flatten())
    #hidden layers
    Model.add(Dense(hidden_Size,activation='tanh'))
    #Dropout layer
    Model.add(Dropout(hidden_dropout))
    #Output layer
    Model.add(Dense(99, activation='softmax'))

    #I am try using different optimizers such as SGD, Adam, RMSProp and I found Adam was the best
    opt1 = keras.optimizers.SGD(learning_rate=learning_rate)
    opt2 = keras.optimizers.Adam(learning_rate=learning_rate)
    opt3 = keras.optimizers.RMSprop(learning_rate=learning_rate)


    Model.compile(loss='sparse_categorical_crossentropy',optimizer=opt2, metrics=['accuracy'] )
    history = Model.fit(x_train, y_train, epochs=50,validation_data=(x_test,y_test))
    acc = history.history['accuracy']
    loss = history.history['loss']
    Model.build(x_train.shape)
    print(Model.summary())
    plt.subplot(1, 2, 1)
    plt.plot(acc)
    plt.subplot(1, 2, 1)
    plt.plot(loss)
    plt.title('model loss & accuracy')
    plt.show()
    print(Model.evaluate(x_test, y_test, verbose=1))
```

I create function that take (batch size, hidden dropout, Regularization, Learning rate, hidden Size, no. epoches)

I call the Function with different values:

ANN(32,0.1,0,0.01,20,100)

```
24/24 [=====] - 0s 4ms/step - loss: 4.4353 - accuracy: 0.0391 - val_loss: 4.4647 - val_accuracy: 0.0323
Epoch 7/50
24/24 [=====] - 0s 5ms/step - loss: 4.3972 - accuracy: 0.0512 - val_loss: 4.4317 - val_accuracy: 0.0363
Epoch 8/50
24/24 [=====] - 0s 4ms/step - loss: 4.3394 - accuracy: 0.0553 - val_loss: 4.4017 - val_accuracy: 0.0444
Epoch 9/50
24/24 [=====] - 0s 4ms/step - loss: 4.3127 - accuracy: 0.0660 - val_loss: 4.3714 - val_accuracy: 0.0444
Epoch 10/50
24/24 [=====] - 0s 4ms/step - loss: 4.2822 - accuracy: 0.0782 - val_loss: 4.3435 - val_accuracy: 0.0605
Epoch 11/50
24/24 [=====] - 0s 4ms/step - loss: 4.2409 - accuracy: 0.0930 - val_loss: 4.3134 - val_accuracy: 0.0645
Epoch 12/50
24/24 [=====] - 0s 4ms/step - loss: 4.2119 - accuracy: 0.1092 - val_loss: 4.2868 - val_accuracy: 0.0766
Epoch 13/50
24/24 [=====] - 0s 4ms/step - loss: 4.1705 - accuracy: 0.1186 - val_loss: 4.2606 - val_accuracy: 0.0847
Epoch 14/50
24/24 [=====] - 0s 4ms/step - loss: 4.1325 - accuracy: 0.1429 - val_loss: 4.2343 - val_accuracy: 0.0968
Epoch 15/50
24/24 [=====] - 0s 4ms/step - loss: 4.1003 - accuracy: 0.1496 - val_loss: 4.2087 - val_accuracy: 0.1048
Epoch 16/50
24/24 [=====] - 0s 4ms/step - loss: 4.0686 - accuracy: 0.1577 - val_loss: 4.1861 - val_accuracy: 0.1048
Epoch 17/50
24/24 [=====] - 0s 3ms/step - loss: 4.0313 - accuracy: 0.1739 - val_loss: 4.1630 - val_accuracy: 0.1169
Epoch 18/50
24/24 [=====] - 0s 5ms/step - loss: 4.0033 - accuracy: 0.1671 - val_loss: 4.1388 - val_accuracy: 0.1210
Epoch 19/50
24/24 [=====] - 0s 3ms/step - loss: 3.9788 - accuracy: 0.2116 - val_loss: 4.1159 - val_accuracy: 0.1411
Epoch 20/50
24/24 [=====] - 0s 3ms/step - loss: 3.9448 - accuracy: 0.2102 - val_loss: 4.0925 - val_accuracy: 0.1492
Epoch 21/50
24/24 [=====] - 0s 3ms/step - loss: 3.9225 - accuracy: 0.2291 - val_loss: 4.0702 - val_accuracy: 0.1694
Epoch 22/50
24/24 [=====] - 0s 4ms/step - loss: 3.8898 - accuracy: 0.2534 - val_loss: 4.0482 - val_accuracy: 0.1815
Epoch 23/50
24/24 [=====] - 0s 3ms/step - loss: 3.8553 - accuracy: 0.2493 - val_loss: 4.0269 - val_accuracy: 0.1895
Epoch 24/50
24/24 [=====] - 0s 3ms/step - loss: 3.8268 - accuracy: 0.2763 - val_loss: 4.0049 - val_accuracy: 0.2056
Epoch 25/50
24/24 [=====] - 0s 3ms/step - loss: 3.8075 - accuracy: 0.2763 - val_loss: 3.9832 - val_accuracy: 0.2097
Epoch 26/50
```




8/8 [=====] - 0s 4ms/step - loss: 0.3682 - accuracy: 0.8831  
[0.36819079518318176, 0.8830645084381104]

➤ I tried different values of batch sizes

```
#try different values of batch sizes  
batch_sizes=[64,128,256]  
for i in batch_sizes:  
    ANN(i,0.1,0,0.01,20,100)
```


I found this the best model (batch\_sizes=128)



➤ I tried different values of hidden sizes

```
hidden_Size=[10,20,30,40]
for i in hidden_Size:
    ANN(64,0.1,0,0.01,i,100)
```

I found this the best model (hidden\_size=40)






➤ I tried different values of hidden dropout

```
55] hidden_dropout=[0,0.1,0.3,0.5,0.7]
    for i in hidden_dropout:
        ANN(64,i,0,0.01,20,100)
```


I found this the best model (hidden\_dropout=0.1)



➤ I tried different values of learning rate


```
#try different values of learning_rates
learning_rates=[0.1,0.03,0.01,0.0003]
for i in learning_rates:
    ANN(64,0.1,0,i,20,100)
```

I found this the best model (learning\_rates=0.01)




- When training a model, it is often useful to lower the learning rate as the training progresses. This schedule applies an exponential decay function to an optimizer step, given a provided initial learning rate.

```
learningRate_schedule = keras.optimizers.schedules.ExponentialDecay(  
    initial_learning_rate=1e-2,  
    decay_steps=10000,  
    decay_rate=0.9)  
ANN(64,0.1,0,learningRate_schedule,20,100)
```



```
learningRate_schedule = keras.optimizers.schedules.ExponentialDecay(  
    initial_learning_rate=0.02,  
    decay_steps=10000,  
    decay_rate=0.9)  
ANN(64,0.1,0,learningRate_schedule,20,100)
```



## Conclusions:

- **I am try using different optimizers such as SGD, Adam, RMSProp and I found Adam was the best**
- **Best Hyperparameters**  
batch\_sizes=128  
hidden\_size=40  
hidden\_dropout=0.1  
learning\_rates=0.01

**Best accuracy is 96.37%**

## References

- <https://www.kaggle.com/c/leaf-classification/data>
- [https://www.tensorflow.org/api\\_docs/python/tf/keras/optimizers/schedules/ExponentialDecay](https://www.tensorflow.org/api_docs/python/tf/keras/optimizers/schedules/ExponentialDecay)
- <https://keras.io/api/optimizers/>
- [https://www.tensorflow.org/tutorials/keras/keras\\_tuner](https://www.tensorflow.org/tutorials/keras/keras_tuner)